

## Лабораторная работа №4. Реализация приложения по распознаванию номеров домов.

Данные: Набор изображений из Google Street View с изображениями номеров домов, содержащий 10 классов, соответствующих цифрам от 0 до 9. 73257 изображений цифр в обучающей выборке; 26032 изображения цифр в тестовой выборке; 531131 изображения, которые можно использовать как дополнение к обучающей выборке; В двух форматах: Оригинальные изображения с выделенными цифрами; Изображения размером 32 × 32, содержащих одну цифру; Данные первого формата можно скачать по ссылкам: <http://ufldl.stanford.edu/housenumbers/train.tar.gz> (обучающая выборка); <http://ufldl.stanford.edu/housenumbers/test.tar.gz> (тестовая выборка); <http://ufldl.stanford.edu/housenumbers/extra.tar.gz> (дополнительные данные); Данные второго формата можно скачать по ссылкам: [http://ufldl.stanford.edu/housenumbers/train\\_32x32.mat](http://ufldl.stanford.edu/housenumbers/train_32x32.mat) (обучающая выборка); [http://ufldl.stanford.edu/housenumbers/test\\_32x32.mat](http://ufldl.stanford.edu/housenumbers/test_32x32.mat) (тестовая выборка); [http://ufldl.stanford.edu/housenumbers/extra\\_32x32.mat](http://ufldl.stanford.edu/housenumbers/extra_32x32.mat) (дополнительные данные); Описание данных на английском языке доступно по ссылке: <http://ufldl.stanford.edu/housenumbers/>

### Задание 1.

Реализуйте глубокую нейронную сеть (полносвязную или сверточную) и обучите ее на синтетических данных (например, наборы MNIST (<http://yann.lecun.com/exdb/mnist/>) или potMNIST). Ознакомьтесь с имеющимися работами по данной тематике: англоязычная статья (<http://static.googleusercontent.com/media/research.google.com/en//pubs/archive/42241.pdf>), видео на YouTube ([https://www.youtube.com/watch?v=vGPI\\_JvLoN0](https://www.youtube.com/watch?v=vGPI_JvLoN0)).

In [1]:

```
from tensorflow import keras
import numpy as np
import scipy.io
import os
import h5py
from PIL import Image

max_len = 6
number_of_digits = 10
wildcard = 10

epochs = 5

data_folder = '../data'
```

In [2]:

```
(X_train, y_train), (X_dev, y_dev) = keras.datasets.mnist.load_data()

X_train = X_train.reshape((X_train.shape + (1, )))
X_dev = X_dev.reshape((X_dev.shape + (1, )))

y_train = y_train.reshape((y_train.shape + (1, )))
y_dev = y_dev.reshape((y_dev.shape + (1, )))
```

In [3]:

```
def preprocess(y):
    result = [[] for _ in range(max_len + 1)]
    for element in y:
        digit = str(element[0])

        result[0].append(len(digit))

        for j in range(max_len):
            val = int(digit[j]) if j < len(digit) else wildcard
            result[j + 1].append(val)

    return [np.array(output).reshape((len(y), 1)) for output in result]
```

In [4]:

```
y_train = preprocess(y_train)
```

```
y_dev = preprocess(y_dev)
```

In [5]:

```
def le_net_5_modern_model(input_shape):
    inputs = keras.layers.Input(shape=input_shape)

    hidden1 = keras.layers.Conv2D(6, (3, 3), activation='relu')(inputs)
    hidden1 = keras.layers.MaxPooling2D()(hidden1)

    hidden2 = keras.layers.Conv2D(16, (3, 3), activation='relu')(hidden1)
    hidden2 = keras.layers.MaxPooling2D(padding="same")(hidden2)

    hidden3 = keras.layers.Flatten()(hidden2)
    hidden3 = keras.layers.Dense(120, activation='relu')(hidden3)

    hidden4 = keras.layers.Dense(84, activation='relu')(hidden3)

    max_len_output = keras.layers.Dense(max_len + 1, activation = 'softmax')(hidden4)
    digit_outputs = [keras.layers.Dense(number_of_digits + 1, activation = 'softmax')(hidden4)
                     for _ in range(max_len)]

    return keras.Model(inputs=inputs, outputs=[max_len_output] + digit_outputs)
```

In [6]:

```
def train(model, epochs):
    model.compile(optimizer='adam',
                  loss='sparse_categorical_crossentropy',
                  metrics=['accuracy'])

    model.fit(X_train, y_train, epochs=epochs, batch_size=256)

    return model.evaluate(X_dev, y_dev)
```

In [7]:

```
train(le_net_5_modern_model(X_train.shape[1:]), epochs)
```

Train on 60000 samples

Epoch 1/5

60000/60000 [=====] - 35s 584us/sample - loss: 5.3751 - dense\_2\_loss: 0.0698 - dense\_3\_loss: 2.6253 - dense\_4\_loss: 0.3824 - dense\_5\_loss: 1.0715 - dense\_6\_loss: 0.2605 - dense\_7\_loss: 0.1635 - dense\_8\_loss: 0.7897 - dense\_2\_accuracy: 0.9941 - dense\_3\_accuracy: 0.7749 - dense\_4\_accuracy: 0.9858 - dense\_5\_accuracy: 0.9822 - dense\_6\_accuracy: 0.9891 - dense\_7\_accuracy: 0.9909 - dense\_8\_accuracy: 0.9844

Epoch 2/5

60000/60000 [=====] - 34s 561us/sample - loss: 0.2633 - dense\_2\_loss: 0.0000e+00 - dense\_3\_loss: 0.2630 - dense\_4\_loss: 9.7298e-07 - dense\_5\_loss: 8.6456e-05 - dense\_6\_loss: 3.1257e-05 - dense\_7\_loss: 1.4265e-05 - dense\_8\_loss: 8.6011e-05 - dense\_2\_accuracy: 1.0000 - dense\_3\_accuracy: 0.9359 - dense\_4\_accuracy: 1.0000 - dense\_5\_accuracy: 1.0000 - dense\_6\_accuracy: 1.0000 - dense\_7\_accuracy: 1.0000 - dense\_8\_accuracy: 1.0000

Epoch 3/5

60000/60000 [=====] - 32s 528us/sample - loss: 0.1600 - dense\_2\_loss: 0.0000e+00 - dense\_3\_loss: 0.1596 - dense\_4\_loss: 2.7721e-06 - dense\_5\_loss: 5.3686e-05 - dense\_6\_loss: 2.5267e-05 - dense\_7\_loss: 3.3568e-05 - dense\_8\_loss: 3.1583e-05 - dense\_2\_accuracy: 1.0000 - dense\_3\_accuracy: 0.9570 - dense\_4\_accuracy: 1.0000 - dense\_5\_accuracy: 1.0000 - dense\_6\_accuracy: 1.0000 - dense\_7\_accuracy: 1.0000 - dense\_8\_accuracy: 1.0000

Epoch 4/5

60000/60000 [=====] - 32s 536us/sample - loss: 0.1154 - dense\_2\_loss: 3.9631e-12 - dense\_3\_loss: 0.1152 - dense\_4\_loss: 1.9085e-05 - dense\_5\_loss: 3.2637e-05 - dense\_6\_loss: 2.3285e-05 - dense\_7\_loss: 3.6202e-05 - dense\_8\_loss: 1.3392e-05 - dense\_2\_accuracy: 1.0000 - dense\_3\_accuracy: 0.9682 - dense\_4\_accuracy: 1.0000 - dense\_5\_accuracy: 1.0000 - dense\_6\_accuracy: 1.0000 - dense\_7\_accuracy: 1.0000 - dense\_8\_accuracy: 1.0000

Epoch 5/5

60000/60000 [=====] - 32s 538us/sample - loss: 0.0902 - dense\_2\_loss: 1.1889e-11 - dense\_3\_loss: 0.0901 - dense\_4\_loss: 7.5113e-06 - dense\_5\_loss: 4.3943e-05 - dense\_6\_loss: 2.2469e-05 - dense\_7\_loss: 5.2991e-05 - dense\_8\_loss: 7.6469e-06 - dense\_2\_accuracy: 1.0000 - dense\_3\_accuracy: 0.9740 - dense\_4\_accuracy: 1.0000 - dense\_5\_accuracy: 1.0000 - dense\_6\_accuracy: 1.0000 - dense\_7\_accuracy: 1.0000 - dense\_8\_accuracy: 1.0000  
10000/10000 [=====] - 4s 356us/sample - loss: 0.1127 - dense\_2\_loss: 1.0712e-10 - dense\_3\_loss: 0.1124 - dense\_4\_loss: 3.5252e-06 - dense\_5\_loss: 4.2240e-05 - dense\_6\_loss: 3.9821e-05 - dense\_7\_loss: 6.1822e-06 - dense\_8\_loss: 3.0734e-07 - dense\_2\_accuracy:

```
1.0000 - dense_3_accuracy: 0.9730 - dense_4_accuracy: 1.0000 - dense_5_accuracy: 1.0000 -  
dense_6_accuracy: 1.0000 - dense_7_accuracy: 1.0000 - dense_8_accuracy: 1.0000
```

Out [7]:

```
[0.11270669284309552,  
 1.0711692e-10,  
 0.11243874,  
 3.5251649e-06,  
 4.2239775e-05,  
 3.9821392e-05,  
 6.1821756e-06,  
 3.0734265e-07,  
 1.0,  
 0.973,  
 1.0,  
 1.0,  
 1.0,  
 1.0,  
 1.0]
```

## Задание 2.

После уточнения модели на синтетических данных попробуйте обучить ее на реальных данных (набор Google Street View).  
Что изменилось в модели?

In [8]:

```
training_dataset = scipy.io.loadmat(os.path.join(data_folder, 'train_32x32.mat'))  
X_train, y_train = training_dataset["X"], training_dataset["y"]  
  
X_train = np.moveaxis(X_train, -1, 0)  
  
dev_dataset = scipy.io.loadmat(os.path.join(data_folder, 'test_32x32.mat'))  
X_dev, y_dev = dev_dataset["X"], dev_dataset["y"]  
  
X_dev = np.moveaxis(X_dev, -1, 0)
```

In [9]:

```
y_train[y_train == 10] = 0  
y_dev[y_dev == 10] = 0  
  
y_train = preprocess(y_train)  
y_dev = preprocess(y_dev)
```

In [10]:

```
train(1e_net_5_modern_model(X_train.shape[1:]), epochs)
```

Train on 73257 samples

Epoch 1/5

```
73257/73257 [=====] - 45s 609us/sample - loss: 3.6116 - dense_11_loss: 0.  
2022 - dense_12_loss: 2.6632 - dense_13_loss: 0.1605 - dense_14_loss: 0.2222 - dense_15_loss: 0.13  
30 - dense_16_loss: 0.1603 - dense_17_loss: 0.0644 - dense_11_accuracy: 0.9856 -  
dense_12_accuracy: 0.3099 - dense_13_accuracy: 0.9868 - dense_14_accuracy: 0.9862 -  
dense_15_accuracy: 0.9948 - dense_16_accuracy: 0.9860 - dense_17_accuracy: 0.9900
```

Epoch 2/5

```
73257/73257 [=====] - 43s 593us/sample - loss: 1.0601 - dense_11_loss: 2.  
2654e-04 - dense_12_loss: 1.0586 - dense_13_loss: 4.4827e-04 - dense_14_loss: 1.7930e-04 - dense_1  
5_loss: 5.0013e-05 - dense_16_loss: 1.8471e-04 - dense_17_loss: 2.9696e-04 - dense_11_accuracy: 1.  
0000 - dense_12_accuracy: 0.6764 - dense_13_accuracy: 1.0000 - dense_14_accuracy: 1.0000 -  
dense_15_accuracy: 1.0000 - dense_16_accuracy: 1.0000 - dense_17_accuracy: 1.0000
```

Epoch 3/5

```
73257/73257 [=====] - 41s 562us/sample - loss: 0.7572 - dense_11_loss: 9.  
8369e-05 - dense_12_loss: 0.7567 - dense_13_loss: 2.2222e-04 - dense_14_loss: 1.2206e-04 - dense_1  
5_loss: 2.5367e-05 - dense_16_loss: 1.1776e-04 - dense_17_loss: 1.7715e-04 - dense_11_accuracy: 1.  
0000 - dense_12_accuracy: 0.7783 - dense_13_accuracy: 1.0000 - dense_14_accuracy: 1.0000 -  
dense_15_accuracy: 1.0000 - dense_16_accuracy: 1.0000 - dense_17_accuracy: 1.0000
```

Epoch 4/5

```
73257/73257 [=====] - 41s 561us/sample - loss: 0.6482 - dense_11_loss: 7.  
4954e-05 - dense_12_loss: 0.6472 - dense_13_loss: 1.8383e-04 - dense_14_loss: 9.9375e-05 - dense_1
```

```

5_loss: 3.0684e-05 - dense_16_loss: 7.1611e-05 - dense_17_loss: 1.0096e-04 - dense_11_accuracy: 1.0000 - dense_12_accuracy: 0.8129 - dense_13_accuracy: 1.0000 - dense_14_accuracy: 1.0000 - dense_15_accuracy: 1.0000 - dense_16_accuracy: 1.0000 - dense_17_accuracy: 1.0000
Epoch 5/5
73257/73257 [=====] - 41s 566us/sample - loss: 0.5944 - dense_11_loss: 8.5858e-05 - dense_12_loss: 0.5942 - dense_13_loss: 1.6579e-04 - dense_14_loss: 1.0570e-04 - dense_15_loss: 3.2646e-05 - dense_16_loss: 6.0651e-05 - dense_17_loss: 9.2520e-05 - dense_11_accuracy: 1.0000 - dense_12_accuracy: 0.8287 - dense_13_accuracy: 1.0000 - dense_14_accuracy: 1.0000 - dense_15_accuracy: 1.0000 - dense_16_accuracy: 1.0000 - dense_17_accuracy: 1.0000
26032/26032 [=====] - 9s 355us/sample - loss: 0.7215 - dense_11_loss: 1.0190e-04 - dense_12_loss: 0.7209 - dense_13_loss: 7.2816e-05 - dense_14_loss: 7.1307e-05 - dense_15_loss: 2.0465e-05 - dense_16_loss: 4.2302e-05 - dense_17_loss: 9.6373e-05 - dense_11_accuracy: 1.0000 - dense_12_accuracy: 0.8019 - dense_13_accuracy: 1.0000 - dense_14_accuracy: 1.0000 - dense_15_accuracy: 1.0000 - dense_16_accuracy: 1.0000 - dense_17_accuracy: 1.0000

```

Out [10]:

```

[0.7214683724662315,
 0.00010190105,
 0.72090775,
 7.2816394e-05,
 7.1306946e-05,
 2.0465208e-05,
 4.230245e-05,
 9.6373355e-05,
 1.0,
 0.8019361,
 1.0,
 1.0,
 1.0,
 1.0,
 1.0]

```

In [ ]:

```

def parse_digit_struct(path):
    with h5py.File(path, 'r') as f:
        digit_struct = f['digitStruct']
        bboxes = digit_struct['bbox']
        names = digit_struct['name']

        y = []

        for i in range(len(bboxes)):
            bbox = bboxes[i].item()
            label = f[bbox]['label']

            y_i = [int(label[0][0])] if len(label) <= 1 else [int(f[label[j].item()][0][0]) for j in range(len(label))]

            if len(y_i) > max_len:
                print('Increase max len up to {}'.format(len(y_i)))
                continue

            y_i = int(''.join(map(str, [0 if digit == 10 else digit for digit in y_i])))

            y.append(y_i)

        return np.array(y).reshape((len(y), 1))

```

In [12]:

```

y_train = parse_digit_struct(os.path.join(data_folder, 'train', 'digitStruct.mat'))
y_dev = parse_digit_struct(os.path.join(data_folder, 'test', 'digitStruct.mat'))

y_train = preprocess(y_train)
y_dev = preprocess(y_dev)

```

In [13]:

```

def load_data(name, h, w):
    x = []

    path = os.path.join(data_folder, name)

```

```

files = sorted(
    filter(
        lambda file_name: file_name.endswith('.png'),
        os.listdir(path)),
    key=lambda file_name: int(file_name[: -len('.png')]))

for file_name in files:
    try:
        img_path = os.path.join(path, file_name)
        img = Image.open(img_path)
        img = img.resize((h, w))
        img = np.array(img)

        X.append(img)
    except:
        pass

return np.array(X)

```

In [14]:

```

X_train = load_data('train', 128, 128)
X_dev = load_data('test', 128, 128)

```

In [15]:

```
train(1e_net_5_modern_model(X_train.shape[1:]), epochs)
```

Train on 33402 samples

Epoch 1/5

```

33402/33402 [=====] - 99s 3ms/sample - loss: 19.1338 - dense_20_loss: 3.5
059 - dense_21_loss: 5.0858 - dense_22_loss: 5.1958 - dense_23_loss: 2.7071 - dense_24_loss: 0.940
8 - dense_25_loss: 0.7455 - dense_26_loss: 0.9054 - dense_20_accuracy: 0.5045 - dense_21_accuracy:
0.2287 - dense_22_accuracy: 0.1359 - dense_23_accuracy: 0.6524 - dense_24_accuracy: 0.9422 -
dense_25_accuracy: 0.9837 - dense_26_accuracy: 0.9690

```

Epoch 2/5

```

33402/33402 [=====] - 97s 3ms/sample - loss: 6.7327 - dense_20_loss: 0.91
94 - dense_21_loss: 1.9993 - dense_22_loss: 2.2835 - dense_23_loss: 1.2363 - dense_24_loss: 0.2807
- dense_25_loss: 0.0052 - dense_26_loss: 0.0053 - dense_20_accuracy: 0.6205 - dense_21_accuracy: 0
.3014 - dense_22_accuracy: 0.2168 - dense_23_accuracy: 0.6954 - dense_24_accuracy: 0.9568 -
dense_25_accuracy: 0.9997 - dense_26_accuracy: 1.0000

```

Epoch 3/5

```

33402/33402 [=====] - 98s 3ms/sample - loss: 5.9529 - dense_20_loss: 0.72
35 - dense_21_loss: 1.8027 - dense_22_loss: 2.0580 - dense_23_loss: 1.1218 - dense_24_loss: 0.2402
- dense_25_loss: 0.0040 - dense_26_loss: 0.0024 - dense_20_accuracy: 0.7097 - dense_21_accuracy: 0
.3909 - dense_22_accuracy: 0.3103 - dense_23_accuracy: 0.6967 - dense_24_accuracy: 0.9567 -
dense_25_accuracy: 0.9997 - dense_26_accuracy: 1.0000

```

Epoch 4/5

```

33402/33402 [=====] - 98s 3ms/sample - loss: 5.3776 - dense_20_loss: 0.63
31 - dense_21_loss: 1.6157 - dense_22_loss: 1.8602 - dense_23_loss: 1.0437 - dense_24_loss: 0.2197
- dense_25_loss: 0.0036 - dense_26_loss: 0.0014 - dense_20_accuracy: 0.7453 - dense_21_accuracy: 0
.4587 - dense_22_accuracy: 0.3827 - dense_23_accuracy: 0.7058 - dense_24_accuracy: 0.9567 -
dense_25_accuracy: 0.9997 - dense_26_accuracy: 1.0000

```

Epoch 5/5

```

33402/33402 [=====] - 99s 3ms/sample - loss: 4.9021 - dense_20_loss: 0.56
24 - dense_21_loss: 1.4673 - dense_22_loss: 1.7005 - dense_23_loss: 0.9661 - dense_24_loss: 0.2022
- dense_25_loss: 0.0031 - dense_26_loss: 8.7435e-04 - dense_20_accuracy: 0.7782 -
dense_21_accuracy: 0.5111 - dense_22_accuracy: 0.4430 - dense_23_accuracy: 0.7177 -
dense_24_accuracy: 0.9571 - dense_25_accuracy: 0.9997 - dense_26_accuracy: 1.0000
13068/13068 [=====] - 32s 2ms/sample - loss: 7.3537 - dense_20_loss: 1.22
30 - dense_21_loss: 2.2428 - dense_22_loss: 2.6212 - dense_23_loss: 1.1236 - dense_24_loss: 0.1419
- dense_25_loss: 0.0031 - dense_26_loss: 7.3890e-04 - dense_20_accuracy: 0.5716 -
dense_21_accuracy: 0.3401 - dense_22_accuracy: 0.2533 - dense_23_accuracy: 0.7303 -
dense_24_accuracy: 0.9861 - dense_25_accuracy: 0.9999 - dense_26_accuracy: 1.0000

```

Out [15]:

```

[7.35365374576509,
 1.2230496,
 2.242823,
 2.6211615,
 1.1236367,
 0.14186539,
 0.0031170375.]

```

0.0001170010,  
0.000738905,  
0.57162535,  
0.34014386,  
0.25329047,  
0.7303336,  
0.98607284,  
0.99992347,  
1.0]